

newbook-fromscratch

February 29, 2024

So, here's the story of the people who are travelling to a destination from various HomePlanets. But sadly they met with a unexpected accident and some people got transported from the ship. So our job is to find if a person with some details can be transported or stayed in ship.

Let's not wait any minute,

We need some packages to be imported and they are imported here, some may be imported later in the notebook.

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Our data is brought into the book by these statements.

```
[3]: og_train_data = pd.read_csv('train.csv')
og_test_data = pd.read_csv('test.csv')

data = og_train_data
```

Let's take a peek at the data.

```
[4]: data.head()
```

```
[4]: PassengerId HomePlanet CryoSleep Cabin Destination Age VIP \
0      0001_01      Europa      False B/0/P TRAPPIST-1e 39.0 False
1      0002_01       Earth      False F/0/S TRAPPIST-1e 24.0 False
2      0003_01      Europa      False A/0/S TRAPPIST-1e 58.0  True
3      0003_02      Europa      False A/0/S TRAPPIST-1e 33.0 False
4      0004_01       Earth      False F/1/S TRAPPIST-1e 16.0 False

      RoomService FoodCourt ShoppingMall Spa VRDeck      Name \
0           0.0         0.0           0.0  0.0   0.0  Maham Ofracculy
1        109.0          9.0          25.0 549.0  44.0   Juanna Vines
2         43.0       3576.0           0.0 6715.0  49.0   Altark Susent
3           0.0       1283.0         371.0 3329.0 193.0   Solam Susent
4        303.0         70.0         151.0  565.0   2.0  Willy Santantines
```

Transported

```

0      False
1       True
2      False
3      False
4       True

```

```
[5]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8693 entries, 0 to 8692
Data columns (total 14 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   PassengerId     8693 non-null   object
 1   HomePlanet      8492 non-null   object
 2   CryoSleep       8476 non-null   object
 3   Cabin           8494 non-null   object
 4   Destination     8511 non-null   object
 5   Age             8514 non-null   float64
 6   VIP             8490 non-null   object
 7   RoomService     8512 non-null   float64
 8   FoodCourt       8510 non-null   float64
 9   ShoppingMall    8485 non-null   float64
10   Spa             8510 non-null   float64
11   VRDeck          8505 non-null   float64
12   Name            8493 non-null   object
13   Transported     8693 non-null   bool
dtypes: bool(1), float64(6), object(7)
memory usage: 891.5+ KB

```

The data types of the attributes are as follows.

```
[6]: data.dtypes
```

```

[6]: PassengerId     object
     HomePlanet      object
     CryoSleep       object
     Cabin           object
     Destination     object
     Age             float64
     VIP             object
     RoomService     float64
     FoodCourt       float64
     ShoppingMall    float64
     Spa             float64
     VRDeck          float64
     Name            object
     Transported     bool

```

dtype: object

Are there any duplicated rows in the data?

```
[7]: data[data.duplicated()]
```

[7]: Empty DataFrame

Columns: [PassengerId, HomePlanet, CryoSleep, Cabin, Destination, Age, VIP, RoomService, FoodCourt, ShoppingMall, Spa, VRDeck, Name, Transported]
Index: []

- There are no duplicated rows in the data

How does the data can be described?

```
[8]: data.describe()
```

```
[8]:
```

	Age	RoomService	FoodCourt	ShoppingMall	Spa \
count	8514.000000	8512.000000	8510.000000	8485.000000	8510.000000
mean	28.827930	224.687617	458.077203	173.729169	311.138778
std	14.489021	666.717663	1611.489240	604.696458	1136.705535
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	19.000000	0.000000	0.000000	0.000000	0.000000
50%	27.000000	0.000000	0.000000	0.000000	0.000000
75%	38.000000	47.000000	76.000000	27.000000	59.000000
max	79.000000	14327.000000	29813.000000	23492.000000	22408.000000

	VRDeck
count	8505.000000
mean	304.854791
std	1145.717189
min	0.000000
25%	0.000000
50%	0.000000
75%	46.000000
max	24133.000000

Are there any missing values in the data? If so, then how many?

```
[9]: data.isnull().sum()
```

```
[9]: PassengerId      0
     HomePlanet     201
     CryoSleep      217
     Cabin          199
     Destination    182
     Age            179
     VIP            203
     RoomService    181
```

```
FoodCourt      183
ShoppingMall   208
Spa            183
VRDeck         188
Name           200
Transported     0
dtype: int64
```

What is the shape of the data?

```
[10]: data.shape
```

```
[10]: (8693, 14)
```

columns are:

```
[11]: data.columns
```

```
[11]: Index(['PassengerId', 'HomePlanet', 'CryoSleep', 'Cabin', 'Destination', 'Age',
            'VIP', 'RoomService', 'FoodCourt', 'ShoppingMall', 'Spa', 'VRDeck',
            'Name', 'Transported'],
            dtype='object')
```

How many unique values are there for the features?

```
[12]: data.nunique()
```

```
[12]: PassengerId      8693
HomePlanet           3
CryoSleep            2
Cabin               6560
Destination          3
Age                  80
VIP                  2
RoomService         1273
FoodCourt           1507
ShoppingMall        1115
Spa                 1327
VRDeck              1306
Name                8473
Transported          2
dtype: int64
```

- I think we can say that the transportation doesn't depend on the 'PassengerId' and 'Name' of the people.
- So, I am going to remove them from the data and store it in a new DataFrame.

```
[13]: lowDim_data = data.drop(columns = ['PassengerId', 'Name'])
```

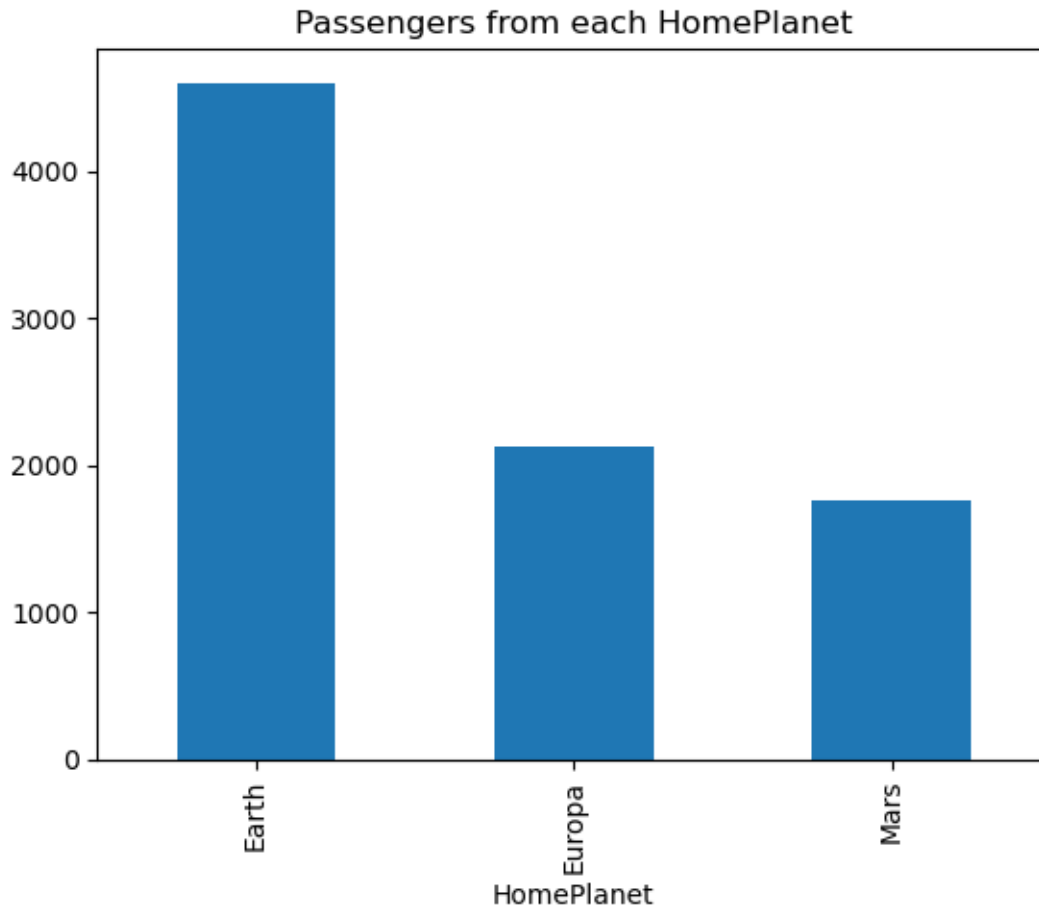
```
[14]: lowDim_data.head()
```

```
[14]:  HomePlanet CryoSleep  Cabin  Destination  Age  VIP  RoomService  \
0      Europa      False B/O/P  TRAPPIST-1e  39.0 False         0.0
1       Earth      False F/O/S  TRAPPIST-1e  24.0 False        109.0
2      Europa      False A/O/S  TRAPPIST-1e  58.0  True         43.0
3      Europa      False A/O/S  TRAPPIST-1e  33.0 False         0.0
4       Earth      False F/1/S  TRAPPIST-1e  16.0 False        303.0

      FoodCourt  ShoppingMall      Spa  VRDeck  Transported
0          0.0          0.0    0.0    0.0         False
1          9.0         25.0  549.0   44.0          True
2        3576.0          0.0  6715.0   49.0         False
3        1283.0        371.0  3329.0  193.0         False
4         70.0        151.0   565.0    2.0          True
```

Let's see that how many are from each HomePlanet

```
[15]: # def addlabels(x,y):
#     for i in range(len(x)):
#         plt.text(i, y[i], y[i], ha = 'center')
ax = data['HomePlanet'].value_counts() \
      .plot(kind = 'bar', title = 'Passengers from each HomePlanet')
# ax.addlabels()
plt.show()
```



```
[16]: data['HomePlanet'].value_counts()
```

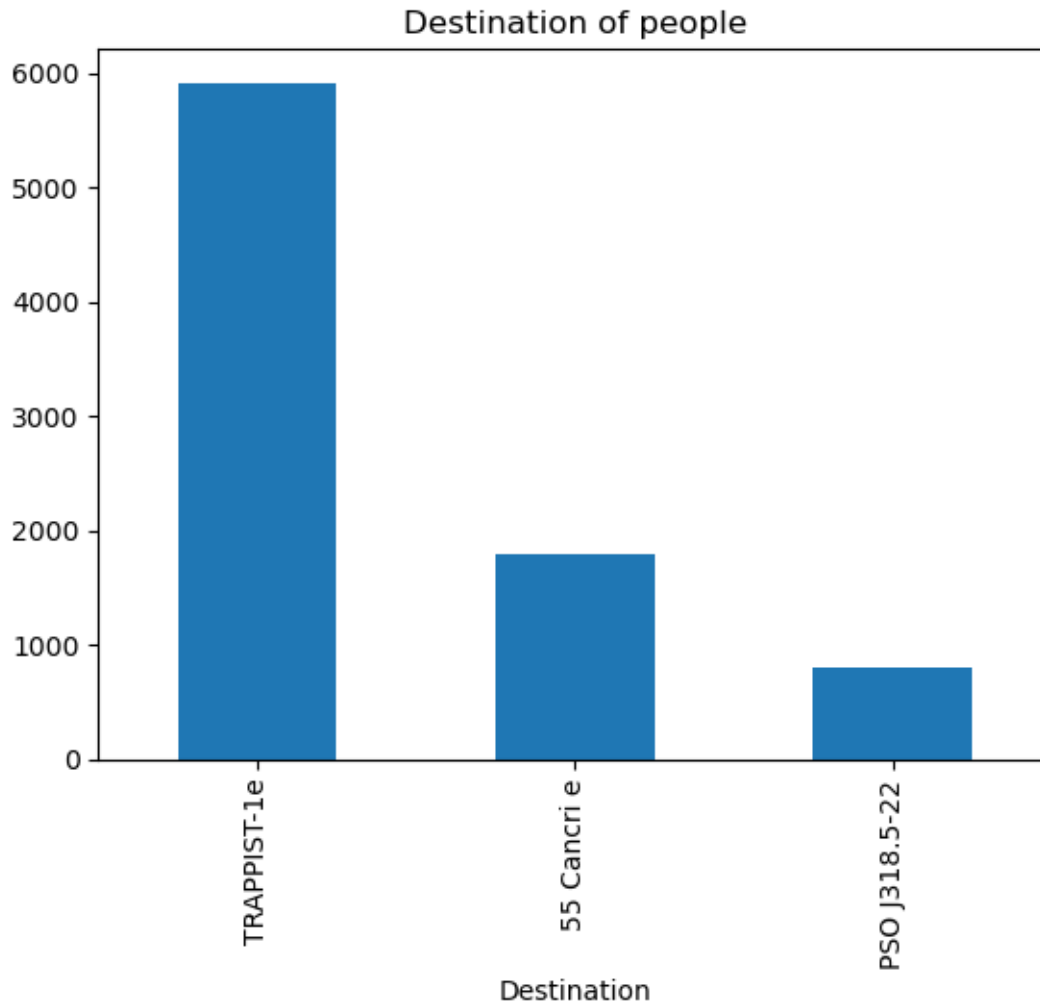
```
[16]: HomePlanet
      Earth      4602
      Europa    2131
      Mars      1759
      Name: count, dtype: int64
```

Maximum people are from Earth. Europa and Mars passengers are close in numbers but far less than the number from earth.

Where do the passengers want to go?

```
[17]: data['Destination'].value_counts().plot(kind = 'bar', title = 'Destination of_
      ↳people')
```

```
[17]: <Axes: title={'center': 'Destination of people'}, xlabel='Destination'>
```



I'm going to fill the Null values of HomePlanet and Destination with Mode of the Cols

```
[18]: data['Destination'].value_counts()
```

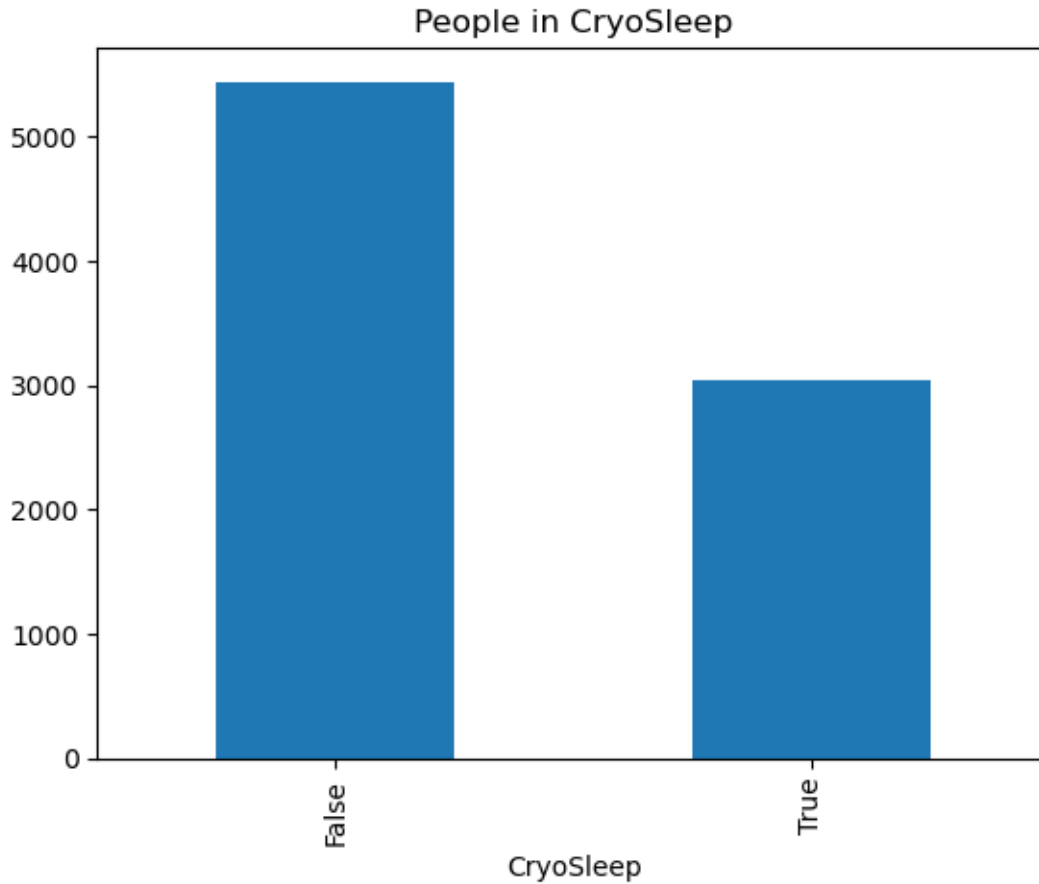
```
[18]: Destination
      TRAPPIST-1e      5915
      55 Cancri e      1800
      PSO J318.5-22      796
      Name: count, dtype: int64
```

- Maximum people are going to the TRAPPIST-1e

```
[19]: lowDim_data['HomePlanet'].fillna('Earth', inplace=True)
      lowDim_data['Destination'].fillna('TRAPPIST-1e', inplace=True)
```

So, how many people are in CryoSleep?

```
[20]: data['CryoSleep'].value_counts().plot(kind = 'bar', title = 'People in_CryoSleep')
plt.figure(figsize = (10, 10))
plt.show()
```



<Figure size 1000x1000 with 0 Axes>


```
[21]: data['CryoSleep'].value_counts()
```

```
[21]: CryoSleep
False    5439
True      3037
Name: count, dtype: int64
```

More people are not in CryoSleep

What is the age distribution in passengers?

```
[22]: data['Age'].plot(kind = 'kde', title = 'Age distribution', x = 'Age')
plt.show()
```

```
[23]: data['Age'].value_counts().head(20)
```

```
[23]: Age
24.0    324
18.0    320
21.0    311
19.0    293
23.0    292
22.0    291
20.0    277
26.0    268
28.0    267
27.0    259
25.0    243
29.0    230
31.0    202
32.0    199
30.0    183
33.0    178
36.0    178
0.0     178
```

```
37.0    177
35.0    171
Name: count, dtype: int64
```

We can observe that most people are of from range (15 - 40) approximately.

How many got transported?

```
[24]: data['Transported'].value_counts()
```

```
[24]: Transported
      True      4378
      False     4315
      Name: count, dtype: int64
```

- Nearly Half of the people got transported.

I'm converting some columns data types to the appropriate types. Because CryoSleep and VIP are actually boolean but in the dataset they are observed as object datatype.

```
[25]: lowDim_data['CryoSleep'] = lowDim_data['CryoSleep'].astype('bool')
      lowDim_data['VIP'] = lowDim_data['VIP'].astype('bool')
```

And also converting HomePlanet, Cabin, Destination to string datatype which are in object datatype.

```
[26]: for i in ['HomePlanet', 'Cabin', 'Destination']:
      lowDim_data[i] = lowDim_data[i].astype('str')
```

- Checking that how many null values are there

```
[27]: lowDim_data.isna().sum()
```

```
[27]: HomePlanet      0
      CryoSleep      0
      Cabin          0
      Destination    0
      Age            179
      VIP            0
      RoomService    181
      FoodCourt       183
      ShoppingMall    208
      Spa             183
      VRDeck          188
      Transported     0
      dtype: int64
```

Now i'm going to fill these null values with their median, cause median gives us the middle value of data which is not effected by the outliers.

Note that i didn't removed the outliers.

```
[28]: for i in ['Age', 'RoomService', 'FoodCourt', 'ShoppingMall', 'Spa', 'VRDeck']:
      lowDim_data[i] = lowDim_data[i].fillna(float(lowDim_data[i].median()))
```

Again checking the data for the Null values.

```
[29]: lowDim_data.isnull().sum()
```

```
[29]: HomePlanet      0
      CryoSleep      0
      Cabin          0
      Destination    0
      Age            0
      VIP            0
      RoomService    0
      FoodCourt      0
      ShoppingMall    0
      Spa            0
      VRDeck         0
      Transported    0
      dtype: int64
```

Now there are no Null values in the data.

Encoding the columns of the dataframe That is i'm converting the string type values to numerical values for the implementation of the model.

```
[30]: from sklearn import preprocessing
      # label_encoder object knows how to understand word labels.
      label_encoder = preprocessing.LabelEncoder()

      # df['species'].unique()
      for i in ['HomePlanet', 'CryoSleep', 'Cabin', 'Destination',
               ↪ 'VIP', 'Transported']:
          lowDim_data[i] = label_encoder.fit_transform(lowDim_data[i])
```

checking the datatypes of the data before model fitting.

```
[31]: lowDim_data.dtypes
```

```
[31]: HomePlanet      int32
      CryoSleep      int64
      Cabin          int32
      Destination    int32
      Age            float64
      VIP            int64
      RoomService    float64
```

```

FoodCourt      float64
ShoppingMall   float64
Spa            float64
VRDeck         float64
Transported    int64
dtype: object

```

This is how our data looks after the encoding

```
[32]: lowDim_data.head()
```

```

[32]:   HomePlanet  CryoSleep  Cabin  Destination  Age  VIP  RoomService  \
0           1           0    149             2  39.0    0           0.0
1           0           0   2184             2  24.0    0          109.0
2           1           0     1             2  58.0    1           43.0
3           1           0     1             2  33.0    0           0.0
4           0           0   2186             2  16.0    0          303.0

      FoodCourt  ShoppingMall      Spa  VRDeck  Transported
0           0.0           0.0     0.0     0.0           0
1           9.0          25.0   549.0    44.0           1
2        3576.0           0.0  6715.0    49.0           0
3        1283.0          371.0  3329.0   193.0           0
4         70.0          151.0   565.0     2.0           1

```

How the variables are correlated?

```
[33]: lowDim_data.corr()
```

```

[33]:   HomePlanet  CryoSleep  Cabin  Destination  Age  \
HomePlanet    1.000000  0.083239 -0.470328    0.034737  0.133577
CryoSleep     0.083239  1.000000  0.112433   -0.094061 -0.073406
Cabin        -0.470328  0.112433  1.000000    0.117281 -0.236994
Destination   0.034737 -0.094061  0.117281    1.000000 -0.006771
Age           0.133577 -0.073406 -0.236994   -0.006771  1.000000
VIP           0.085260 -0.050342 -0.094719   -0.035430  0.058490
RoomService   0.211751 -0.240750 -0.073802    0.045733  0.068629
FoodCourt     0.071454 -0.202675 -0.260056   -0.111057  0.127390
ShoppingMall  0.101383 -0.207213 -0.067709    0.024721  0.033148
Spa           0.055047 -0.196893 -0.180768   -0.055815  0.120946
VRDeck        0.039824 -0.190437 -0.210865   -0.073293  0.099590
Transported   0.115461  0.451744 -0.052604   -0.108152 -0.074233

      VIP  RoomService  FoodCourt  ShoppingMall      Spa  \
HomePlanet  0.085260    0.211751  0.071454    0.101383  0.055047
CryoSleep   -0.050342   -0.240750 -0.202675   -0.207213 -0.196893
Cabin       -0.094719   -0.073802 -0.260056   -0.067709 -0.180768
Destination -0.035430    0.045733 -0.111057    0.024721 -0.055815

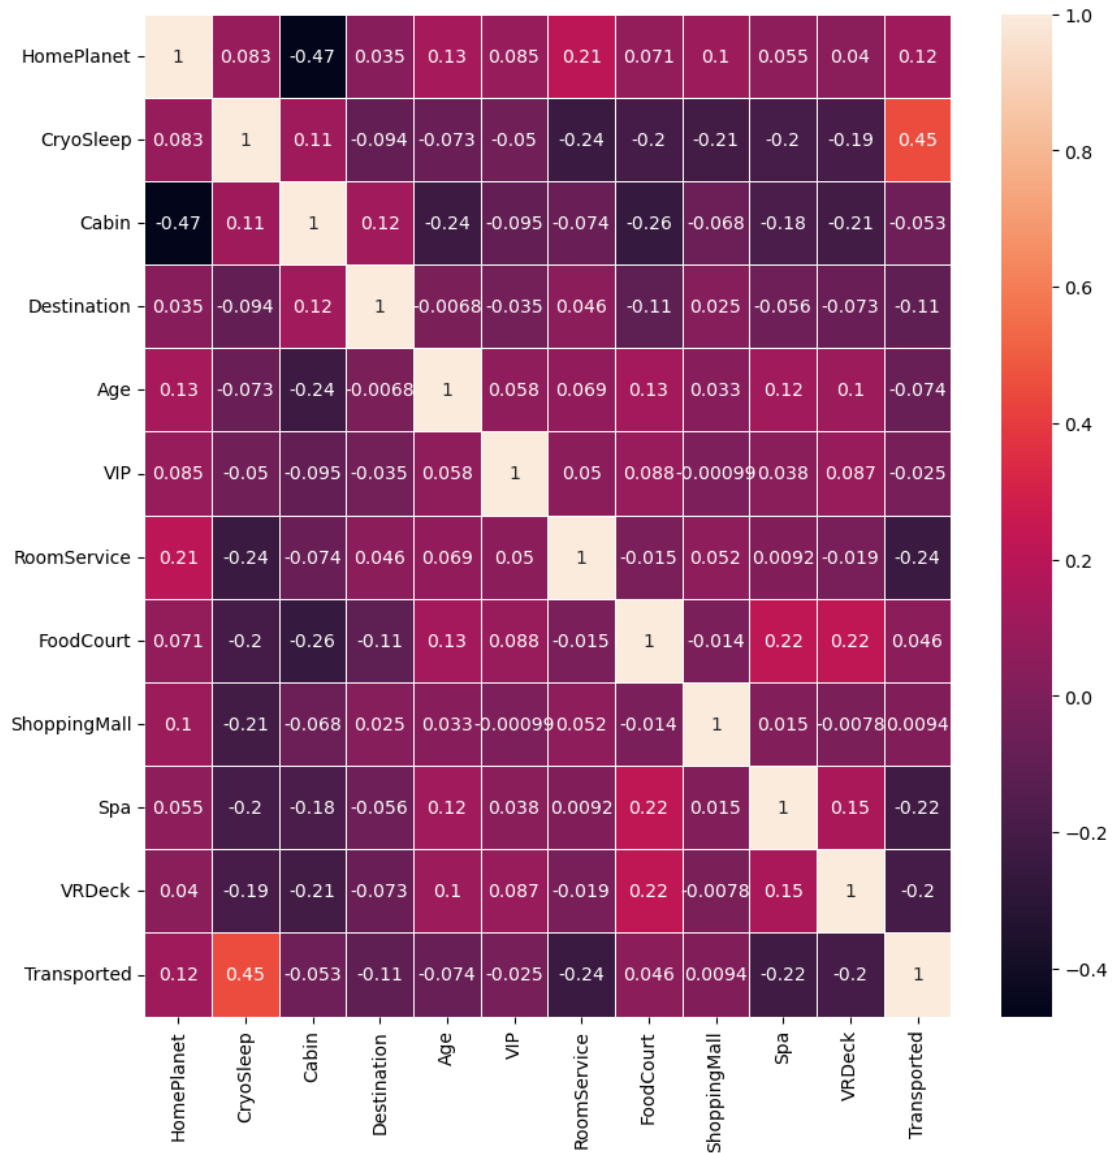
```

Age	0.058490	0.068629	0.127390	0.033148	0.120946
VIP	1.000000	0.050354	0.088208	-0.000992	0.037896
RoomService	0.050354	1.000000	-0.015126	0.052337	0.009244
FoodCourt	0.088208	-0.015126	1.000000	-0.013717	0.221468
ShoppingMall	-0.000992	0.052337	-0.013717	1.000000	0.014542
Spa	0.037896	0.009244	0.221468	0.014542	1.000000
VRDeck	0.087235	-0.018624	0.224572	-0.007849	0.147658
Transported	-0.024602	-0.241124	0.045583	0.009391	-0.218545

	VRDeck	Transported
HomePlanet	0.039824	0.115461
CryoSleep	-0.190437	0.451744
Cabin	-0.210865	-0.052604
Destination	-0.073293	-0.108152
Age	0.099590	-0.074233
VIP	0.087235	-0.024602
RoomService	-0.018624	-0.241124
FoodCourt	0.224572	0.045583
ShoppingMall	-0.007849	0.009391
Spa	0.147658	-0.218545
VRDeck	1.000000	-0.204874
Transported	-0.204874	1.000000

```
[34]: fig, ax = plt.subplots(figsize = (10, 10))
      sns.heatmap(lowDim_data.corr(), annot = True, linewidths=.5, ax=ax)
```

```
[34]: <Axes: >
```



From the correlation plot we can see that there are no seriously correlated variables in the data. The highest value is between Cryosleep and Transported with 0.45

Splitting training data into train and validate datasets

```
[35]: x = lowDim_data.iloc[:, :-1]
      y = lowDim_data.iloc[:, -1]
```

```
[36]: from sklearn.model_selection import train_test_split
      x_train, x_validate, y_train, y_validate = train_test_split(x, y, test_size=0.2)
```

```
[37]: y_validate.head()
```

```
[37]: 6885    1
      925    1
      3861   0
      4631   1
      1558   0
      Name: Transported, dtype: int64
```

We are using KNN classifier for our data, ofcourse there are other models but KNN is also a powerful model.

```
[38]: from sklearn.neighbors import KNeighborsClassifier
      knn = KNeighborsClassifier(n_neighbors=3)
      knn.fit(x_train, y_train)
```

```
[38]: KNeighborsClassifier(n_neighbors=3)
```

Importing the test data

```
[39]: test_data = pd.read_csv('test.csv')
```

```
[40]: test_data.head()
```

```
[40]: PassengerId HomePlanet CryoSleep Cabin Destination Age VIP \
0      0013_01      Earth      True  G/3/S  TRAPPIST-1e  27.0  False
1      0018_01      Earth     False  F/4/S  TRAPPIST-1e  19.0  False
2      0019_01     Europa      True  C/0/S  55 Cancr i e  31.0  False
3      0021_01     Europa     False  C/1/S  TRAPPIST-1e  38.0  False
4      0023_01      Earth     False  F/5/S  TRAPPIST-1e  20.0  False

      RoomService FoodCourt ShoppingMall Spa VRDeck Name
0              0.0        0.0          0.0  0.0   0.0  Nelly Carsoning
1              0.0         9.0          0.0 2823.0  0.0   Lerome Peckers
2              0.0         0.0          0.0   0.0   0.0   Sabih Unhearfus
3              0.0       6652.0          0.0  181.0  585.0  Meratz Caltilter
4             10.0         0.0         635.0   0.0   0.0  Brence Harperez
```

Now we'll do all the steps that are done on the train data for the test data to get the test data to the same state as the train data before fitting the model.

```
[41]: test_data.drop(columns=['PassengerId', 'Name'], inplace = True)
```

```
[42]: test_data.isnull().sum()
```

```
[42]: HomePlanet      87
      CryoSleep      93
      Cabin        100
      Destination   92
      Age           91
```

```
VIP          93
RoomService  82
FoodCourt    106
ShoppingMall 98
Spa          101
VRDeck       80
dtype: int64
```

```
[43]: test_data.dtypes
```

```
[43]: HomePlanet    object
      CryoSleep    object
      Cabin        object
      Destination  object
      Age          float64
      VIP          object
      RoomService  float64
      FoodCourt    float64
      ShoppingMall float64
      Spa          float64
      VRDeck       float64
      dtype: object
```

```
[44]: test_data['CryoSleep'] = test_data['CryoSleep'].astype('bool')
      test_data['VIP'] = test_data['VIP'].astype('bool')
```

```
[45]: for i in ['HomePlanet', 'Cabin', 'Destination']:
      test_data[i] = test_data[i].astype('str')
```

```
[46]: test_data.isnull().sum()
```

```
[46]: HomePlanet    0
      CryoSleep    0
      Cabin        0
      Destination  0
      Age          91
      VIP          0
      RoomService  82
      FoodCourt    106
      ShoppingMall 98
      Spa          101
      VRDeck       80
      dtype: int64
```

```
[47]: for i in ['Age', 'RoomService', 'FoodCourt', 'ShoppingMall', 'Spa', 'VRDeck']:
      test_data[i] = test_data[i].fillna(float(test_data[i].mode()))
```



```
C:\Users\91970\AppData\Local\Temp\ipykernel_27200\3196287623.py:2:
FutureWarning: Calling float on a single element Series is deprecated and will
raise a TypeError in the future. Use float(ser.iloc[0]) instead
    test_data[i] = test_data[i].fillna(float(test_data[i].mode()))
```

```
[48]: test_data.isnull().sum()
```

```
[48]: HomePlanet      0
      CryoSleep     0
      Cabin         0
      Destination   0
      Age           0
      VIP           0
      RoomService   0
      FoodCourt     0
      ShoppingMall  0
      Spa           0
      VRDeck        0
      dtype: int64
```

```
[49]: # from sklearn import preprocessing
      # # label_encoder object knows how to understand word labels.
      # label_encoder = preprocessing.LabelEncoder()

      for i in ['HomePlanet', 'CryoSleep', 'Cabin', 'Destination', 'VIP']:
          test_data[i] = label_encoder.fit_transform(test_data[i])
```

```
[50]: test_data.dtypes
```

```
[50]: HomePlanet      int32
      CryoSleep    int64
      Cabin        int32
      Destination  int32
      Age          float64
      VIP          int64
      RoomService  float64
      FoodCourt    float64
      ShoppingMall float64
      Spa          float64
      VRDeck       float64
      dtype: object
```

```
[51]: test_data.head()
```

```
[51]:   HomePlanet  CryoSleep  Cabin  Destination  Age  VIP  RoomService  \
0           0           0       1       2784         2  27.0       0       0.0
```

1	0	0	1867	2	19.0	0	0.0
2	1	1	257	0	31.0	0	0.0
3	1	0	259	2	38.0	0	0.0
4	0	0	1940	2	20.0	0	10.0

	FoodCourt	ShoppingMall	Spa	VRDeck
0	0.0	0.0	0.0	0.0
1	9.0	0.0	2823.0	0.0
2	0.0	0.0	0.0	0.0
3	6652.0	0.0	181.0	585.0
4	0.0	635.0	0.0	0.0

This is how our test data looks after the process

Predicting the Transported values for the test data

```
[52]: knn.predict(test_data)
```

```
[52]: array([1, 0, 1, ..., 1, 1, 1], dtype=int64)
```

Since i dont have the transported values for the test data(actual) i'm taking the validation data from the above splitting and calculate the accuracy score from it.

```
[53]: validate_y = knn.predict(x_validate)
```

This is how the values are predicted by our model

```
[54]: validate_y
```

```
[54]: array([1, 1, 0, ..., 1, 1, 1], dtype=int64)
```

Total how many are Transported and not

```
[55]: validate_y = pd.DataFrame(validate_y) #converting the numpy array structure to pandas dataframe
```

```
[56]: validate_y.value_counts() # 1 means transported, 0 means not transported
```

```
[56]: 1    872
      0    867
      Name: count, dtype: int64
```

This is how the actual validation data Transported values are:

```
[57]: y_validate.value_counts() # 1 means transported, 0 means not transported
```

```
[57]: Transported
      0    882
      1    857
```

Name: count, dtype: int64

Formula for accuracy $\text{accuracy} = (\text{number of correct predictions}) / (\text{total number of predictions})$

```
[58]: y_validate = np.array(y_validate)
      validate_y = np.array(validate_y)
```

```
[59]: from sklearn.metrics import accuracy_score
```

```
[60]: print(f"Our model Accuracy is : {accuracy_score(y_validate, validate_y)} which_
      ↪is not so bad, because we didn't removed the outliers and didn't tuned our _
      ↪ hyperparameters")
```

Our model Accuracy is : 0.7429557216791259 which is not so bad, because we didn't removed the outliers and didn't tuned our hyperparameters